

Construction algorithms

We need to start somewhere

- When we first come to an instance we are looking at an **empty page**
 - In TSP, no cities are visited
 - In Knapsack, all the items are lying on the ground
- Initial solution often created by "constructing" one piece / element at a time

Construction algorithms – Minimum Insert Cost


A better greedy algorithm

- "Minimum Insert" is another **greedy algorithm** – but performs slightly better

- At each iteration
 - Find the customer that, when inserted, increases the **cost by the least**
 - Insert it in the position that increases the cost by the least
 - Allows insert **between any pair of customers** – not just at the end

Construction algorithms – Nearest Neighbour

- For example: TSP
 - Starting with an **empty tour**, start in the "home" city
 - Repeat:
 - Add the **city closest** to the last-added city



Construction algorithms – Greedy


- "Nearest Neighbour" is a classic **greedy algorithm**
- Greedy algorithms perform **locally-optimal moves**
- However, they often make **"structural" errors**.

- Greedy algorithms exist for many problems:
 - Knapsack: Insert the item with the best value/weight ratio
 - Scheduling: Schedule the item with the **earliest deadline**

Construction algorithms

- Combine "greedy" and "random"
- Don't always make the very "greediest" choice

- E.g. choose the j^{th} best of a choices
- $\epsilon = n \cdot \text{uniform}(0,1)^k$



Construction algorithms – Randomisation

- At each iteration
 - Choose a **random customer**
 - Insert it in the **position that increases the cost by the least**
 - Allows insert **between any pair of customers** – not just at the end

Conclusion

- Constructing a solution is the **first step** in solving a combinatorial optimisation problem
- Building up a solution one element at a time is often a good place to start
- Greedy heuristics can give half-good solutions, but are often structurally unsound
- Some techniques can give a heuristic a bit of **foresight**, such as Regret

- Next time
 - Now we have our first solution ... can we **improve it**?

Regret

$$\text{regret}(i) = C(\text{insert } i \text{ in } 2^{\text{nd}} \text{ best route}) - C(\text{insert } i \text{ in best route})$$

$$= f(2, i) - f(1, i)$$

Select the **choice i with largest regret**

- VRP: next customer to be served is the one with largest regret

Can be extended to the k -th best routes:

$$k\text{-regret}(i) = \sum_{j=2}^k (f(j, i) - f(1, i))$$


Regret

- And a little **more structure**...
 - Regret works best when some structure is already present
 - Use **Seeds**.
 - Seed 1: The customer with max (dist to depot)
 - Seed 2: The customer with max (dist to Seed 1)
 - Seed 3: The customer with max (dist to Seed 1 + dist to Seed 2)
 - Seed 4: The customer with max (dist to Seed 1 + dist to Seed 2 + dist to Seed 3)
 - Etc

n-queens


n-queens

- Alg 2: Greedy (a.k.a Hill Climbing)
 - Choose the best move / one of the best moves
 - Requires us to evaluate the entire Neighbourhood



Hill Climbing has been compared to climbing Mt Everest in 1953, and while suffering from amnesia...


- Alg 3: First found
 - Randomise neighbourhood evaluation
 - Make first improving move



Computational evidence favours first-found over greedy


n-queens

- Greedy search is **incomplete**
- Alg 4: Randomised Greedy
 - With probability p do greedy/first found move
 - With otherwise do a random move
 - Asymptotically complete



n-queens

- Alg 5: Biased
 - Choose a move with probability (inversely) proportional to score
 - (Twice as likely to choose a '2' move than a '4')



VRP


Local Search in VRP

- More complex operators
 - 1-move
 - 2-move




Other Neighbourhoods for VRP:

- Swap 1-1




Other Neighbourhoods for VRP:

- Swap 1-1




Other Neighbourhoods for VRP:

- Swap tails



VRP specific operators

- 2-opt (3-opt, 4-opt, ...)
- Remove 2 arcs
- Replace with 2 others



Neighbourhood

- Hard constraints create impenetrable **mountain ranges** in solution space
- Soft constraints allow passes through the mountains

Strongly connected: **hard w/ 2-way reverses only**


Weakly optimally connected: **hard w/ some one-way reverses**

- The optimum can be reached from any starting solution

Neighbourhood

- E.g. Map Colouring (k -colouring)
 - Colour a map (graph) so that no two adjacent countries (nodes) are the same colour
 - Either:
 - Use at most k colours
 - Minimize number of colours

Map Colouring




Escaping Plateaux (Shoulders)


- If no downhill (uphill) moves, allow **sideways** moves in hope that algorithm can escape
 - Need to place a limit on the possible number of sideways moves to avoid infinite loops
- For 8-queens
 - Now allow sideways moves with a limit of 100
 - Raises percentage of problem instances solved from 14% to 94%
- However...
 - 21 steps on average for every successful solution
 - 64 for each failure

Escaping local minima II


- Solution 2: **Simulated Annealing**
 - Based on manner in which crystals are formed
 - At high temperatures, molecules move freely
 - At low temperatures, molecules are "stuck"
 - If cooling is slow
 - A low energy, organized crystal lattice formed
 - Minimise energy in crystal \leftrightarrow Minimise objective



Simulated Annealing



Simulated Annealing



Simulated Annealing accepts/rejects solution candidate based on probability

Simulated Annealing

- If candidate solution **reduces the objective**,
 - always accept the change (c.f. first found)
- If candidate solution has higher objective,
 - system parameter T ("Temperature") controls probability of acceptance
 - $P(\text{accept increase of } \Delta \text{ in objective}) = e^{-\Delta/T}$
 - T reduces as method proceeds
 - As $T \rightarrow 0$, only improving moves accepted

Meta-heuristics: Properties (1)

- can address both discrete- and continuous-domain optimisation problems
- are strategies that "guide" the search process
- range from simple local search procedures to complex adaptive learning processes
- efficiently explore the search space to find **good (near-optimal) feasible solutions**
- provide no guarantee of global or local optimality
- are agnostic to the unexplored feasible space (i.e., no "bound" information)
- lack a metric of "goodness" of solution (often stop due to an external time or iteration limit)

Meta-heuristics: Properties (2)

- are not based on some algebraic model (unlike exact methods)
- can be used in conjunction with an exact method
 - E.g., use meta-heuristic to provide upper bounds
 - E.g., use restricted MILP as "local heuristic" (i.e., math-heuristic)
- are usually non-deterministic
- are not problem specific (but their subordinate heuristics can be)
- may use some form of memory to better guide the search

Meta-heuristics: An Incomplete Survey



Meta-heuristics: Overview

- Exhibit **intensification** and **diversification**
 - Need to do both
 - Can be explicitly controlled
- Intensification (Exploitation):**
 - Concentrate search around already found "good" solutions
 - Look harder in a smaller area
- Diversification (Exploration):**
 - Expand the area being looked at
 - Find new (promising) areas to search
 - Includes mechanisms to avoid getting trapped in confined areas of the search space

Tabu Search

Tabu Search

- Tabu (English):** prohibited, disallowed, forbidden
- Tabu (Javan):** forbidden to use due to being sacred and/or of supernatural powers
- Starts with classic Local Search until a local minimum is found
- Choose an objective-increasing move
- Make undoing that move "tabu"
- Place it on a "tabu list"
- Repeat

Aspiration Criteria

- Allowed to keep a solution if it meets certain criteria
- Most common: a new incumbent / best solution is kept

```

s = GenerateRandomSolution()
TabuList = []
while termination conditions not met do
    s = ChooseBestOf(s, TabuList)
    UpdateTabuList(s)
endwhile
    
```

Soft Constraints

General Idea: Move constraints into the objective

- Relax Constraints**
 - Penalise degree of violation
 - Allows Local Search to move through "infeasibility barrier"
 - Opens new areas of search space
 - Maintain both best feasible solution "best" infeasible solution
 - Increase penalty over time to force incumbent back to feasibility
 - Often used with Simulated Annealing or Tabu Search
 - Extensively used in practice, e.g., VRP
- Parameters:**
 - Which constraints to relax
 - Penalty schedule

Guided Local Search

Basic idea:

- Select a **feature** that indicates a poor solution
- Penalise a solution that exhibits that feature
- Start with zero penalty
- Repeat
 - Perform Local Search to minimise original objective + penalties
 - Select elements to penalise
 - Increase penalty on selected elements

Guided Local Search

Local Search

- Do local search with an **updated objective** $h(s) = g(s) + \lambda \sum p_i \cdot I_i(s)$
- Updated objective
 - $h(s)$: Augmented objective
 - $g(s)$: Original objective
 - λ : "normalisation" parameter
 - p_i : Count of times feature i has been penalised
 - $I_i(s)$: Indicator function: 1 if feature i in solution s ; 0 otherwise

Guided Local Search

Select elements to penalise:

- Update penalty of features that maximise Utility
- c_i : Original cost of feature
- $I_i(s)$: Does solution s exhibit feature i
- At each iteration
 - Local Search using augmented objective
 - Select maximum utility features
 - Set $p_i = p_i + 1$ for all selected features
- Penalty increases each iteration
 - Local Search tries harder to eliminate feature
 - Utility decreases the more often you penalise a feature
 - Eventually select other features

Guided Local Search - TSP

$$h(s) = g(s) + \lambda \sum p_i \cdot I_i(s)$$

$$util(s, f) = I_i(s) \cdot \frac{c_i}{1 + p_i}$$

- Feature: an arc (i, j)
- $I_i(s) = 1$ if arc (i, j) is present; 0 otherwise
- At each iteration
 - Penalise the longest arc
 - Try harder to get rid of it
- As the search progresses
 - Utility of first arc decreases
 - Other arcs start being penalised

Diversification:

- Big problem is getting a homogenous population
- Too much intensification, not enough diversification
- Some algorithms explicitly measure diversity
 - keep lower-quality solutions that maintain diversity
- Meta-meta Soft constraints
 - Maintain a separate population of infeasible solutions

Solution Representation is key

- Needs to fulfil multiple goals
 - Easy to calculate fitness (objective)
 - Easy to perform crossover (search)
 - Easy to manipulate (mutation)
 - Easy for local search
- E.g. VRP
 - First attempts used array for each route, or successor info
 - Very difficult for crossover
 - Better rep turns out to be a single array

Ant Colony Optimisation (ACO)

E.g.: Want to find shortest paths in a communications network

- Send out ants that choose a path
 - Partly at random (our naive heuristic)
 - Partly influenced by previous ants: Pheromone trail
- The first ant to get to a destination increases the Pheromone on its path
- Pheromone levels decrease over time
- More ants select the best path
- The best path gets reinforced


(ACO) - TSP

Probability that ant chooses arc ij : $\tau_{ij} \cdot \eta_{ij}$

Path Relinking

Basic idea:

- Take two solutions
- "Walk" between them



Iterated Local Search

Perturbation

- A.k.a. "Shake" or "Stir"
- "Strength" = how many elements to change

Similar to Tabu Search

- Fix too much \rightarrow Fall back to same solution
- Fix too little \rightarrow Same as random restart

Can store some solution history

- Use it to control perturbation

Variable Neighbourhood Search

- E.g. neighbourhoods for VRP
 - 1-move (move 1 visit to another position)
 - 2-1 swap (swap visits in 2 routes)
 - 2-2 swap (swap 2 visits between 2 routes)
 - Or-opt size 2 (move chain of length 2 - forwards and backwards - anywhere)
 - Or-opt size 3 (chain length 3)
 - Tail exchange (swap final portion of routes)
 - 2-opt
 - 3-opt

GRASP

- Greedy Randomised Ascent Procedure**
 - Similar to ILS (independently developed)
 - "Kick" in ILS is replaced by **restricted construction method**
 - Inserts elements of the solution one at a time
 - Greedy and randomises a good order for insertion
 - Uses a **Restricted Candidate List (RCL)** to guide insertion
- Based on "Feature Cost"
 - E.g. TSP: Feature = City
 - Feature cost = Cost of inserting city into solution (Min Insert Cost)

Variable Neighbourhood Search

Basic idea:

- If you have reached the local min for one neighbourhood, swap neighbourhoods (to a larger one)
- Given a list of neighbourhoods $N = \{N_1, N_2, \dots, N_k\}$
 - Start by applying the first
 - When you get to local minimum, move to next
 - When you find an improvement go back to first

Very successful method across multiple domains

Large Neighbourhood Search

Destroy part of the solution (Select method)

- Examples
 - Choose longest (worst) arc in solution
 - Remove visits at each end
 - Remove nearby visits

Actually, choose x^{th} worst


$$x = n \cdot \text{uniform}(0,1)$$

- unif \rightarrow (unif)
- 0.56 \rightarrow 0.016
- 0.96 \rightarrow 0.531

Destroy part of the solution (Select method)

- Examples
 - Choose first visit randomly
 - Then, remove "related" visits

Based on distance, time compatibility, load



Large Neighbourhood Search

Large Neighbourhood Search

Re-create solution

- Systematic search (e.g. MILP, Constrained Programming, etc)
- Smaller problem, easier to solve
- Can be very effective
- Use your favourite insert method
- Better still, use a portfolio of insert methods

Regal's paper: select amongst

- Minimum Insert Cost
- Regret (2-regret)
- 2-regret
- 4-regret
- Random Insert order

Adaptive Insert Method

Regal adapts choice based on prior performance

"Good" methods are chosen more often

Probability of choosing method



Adaptive Select Method (destruction method)

Probability of choosing method

